

## 2. Finding Geographic Locations of Headlines

October 10, 2019

### 1 Finding Geographic Locations of Headlines

#### 1.1 Adding Latitude and Longitude Coordinates

**Objective:** Find the geographic location of each headline in latitude and longitude coordinates from the city/country names. We will use these coordinates to perform clustering of geographically similar headlines in the next part.

**Workflow:** 1. Load in the Pandas DataFrame with headline, countries, and cities. - If a headline contains multiple cities/countries, decide which single one to keep. 2. For each city/country, match the name to the latitude and longitude in geonamescache. - You can use the function `gc.get_cities_by_names("city_name")`. - Some cities will return multiple matches with the previous function in different countries. You'll have to decide which city to keep based on a heuristic (rule of thumb). - If you have trouble, work with a single problematic city until you figure it out, then write a function to apply on all headlines. 3. Add longitude and latitude coordinates to your DataFrame for each headline. - It will be helpful to get the countrycode of each headline at this point. - If you were not able to find many countries, think about dropping the column. You also need to decide what to do with headlines that have no coordinates. - You should end up with over 600 headlines that have geographic coordinates

**Deliverable:**

The deliverable is a Jupyter Notebook documenting your work as you add three additional columns to the DataFrame: longitude, latitude, and countrycode. We will use these coordinates to cluster the headlines in the next section.

#### 1.2 Read Data into a DataFrame

We stored the headline, cities, and countries as a json file that was a list of dictionaries. This can be directly read in a Pandas dataframe.

```
[1]: import pandas as pd
import numpy as np

data = pd.read_json("../data/headline_cities_and_countries.json")
data = data.replace({None: np.nan})
data.head()
```

```
[1]:
```

	headline	countries	cities
0	Zika Outbreak Hits Miami	NaN	Miami

1	Could Zika Reach New York City?	NaN	New York City
2	First Case of Zika in Miami Beach	NaN	Miami Beach
3	Mystery Virus Spreads in Recife, Brazil	Brazil	Recife
4	Dallas man comes down with case of Zika	NaN	Dallas

```
[2]: data.iloc[3:5]
```

```
[2]:
```

	headline	countries	cities
3	Mystery Virus Spreads in Recife, Brazil	Brazil	Recife
4	Dallas man comes down with case of Zika	NaN	Dallas

We'll rename the columns to singular (since they only have one value each).

```
[3]: data = data.rename(columns=dict(countries="country", cities="city"))
data.tail()
```

```
[3]:
```

	headline	country	city
645	Rumors about Rabies spreading in Jerusalem hav...	NaN	Jerusalem
646	More Zika patients reported in Indang	NaN	Indang
647	Suva authorities confirmed the spread of Rotav...	NaN	Suva
648	More Zika patients reported in Bella Vista	NaN	Bella Vista
649	Zika Outbreak in Wichita Falls	NaN	Wichita Falls

From a brief look at some of the headlines and cities, our regular expression pattern matching looks to have worked well. As we go through the project, we'll keep an eye out for places it may have failed.

### 1.3 Investigate the Data

We can start off using the `.describe()` method to understand our data.

```
[4]: data.describe()
```

```
[4]:
```

	headline	country	city
count	650	15	608
unique	647	10	573
top	Barcelona Struck by Spanish Flu	Brazil	Madrid
freq	2	3	4

It looks like there may be some duplicates in the data since at least one headline is mentioned twice. Let's check for duplicates and then drop any that are duplicated.

```
[5]: data["headline"].value_counts().sort_values().tail()
```

```
[5]:
```

Zika Outbreak in Hyderabad	1
Ibadan tests new cure for Malaria	1
Spanish Flu Outbreak in Lisbon	2
Spanish Flu Spreading through Madrid	2

Barcelona Struck by Spanish Flu                      2  
Name: headline, dtype: int64

```
[6]: print(f"There were {len(data)} rows before dropping duplicates.")  
data = data.drop_duplicates()  
print(f"There are {len(data)} rows after dropping duplicates.")
```

There were 650 rows before dropping duplicates.

There are 647 rows after dropping duplicates.

Another useful method for data investigation is `.info()`

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 647 entries, 0 to 649  
Data columns (total 3 columns):  
headline      647 non-null object  
country       15 non-null object  
city          605 non-null object  
dtypes: object(3)  
memory usage: 20.2+ KB
```

We can see there are many missing countries (635) and some missing cities (42). The data types look correct at this point.


### 1.3.1 Exploratory Plots

Plots are a great way to visualize data. Let's take a look at the distribution of countries and cities.

```
[8]: data['country'].value_counts()
```

```
[8]: Brazil      3  
     Malaysia   3  
     Panama     2  
     Vietnam    1  
     Singapore  1  
     Mexico     1  
     Belize     1  
     Thailand   1  
     Guatemala  1  
     Hong Kong  1  
     Name: country, dtype: int64
```

```
[9]: %matplotlib inline  
_ = data['country'].value_counts().plot.bar(title='Count of Countries')
```



We have many more cities, so a bar plot might not be the best graphic.

```
[10]: print(f'There are {data["country"].nunique()} different countries.')
      print(f'There are {data["city"].nunique()} different cities.')
```

There are 10 different countries.

There are 573 different cities.


Let's just look at the 10 most common cities.

```
[11]: data["city"].value_counts().sort_values().tail(10)
```

```
[11]: Hollywood      2
      Jefferson      2
      Palm Springs   2
      London         2
      Panama City     2
      Cancun          2
      Edinburgh       2
      Madrid          3
      Monroe          4
      Miami           4
```


Name: city, dtype: int64

```
[12]: _ = (
    data["city"]
    .value_counts()
    .sort_values()
    .tail(10)
    .plot.bar(title="10 Most Common Cities")
)
```



We can see that there are no cities that dominate the headlines.

```
[13]: _ = data["city"].value_counts().plot.hist(title="Distribution of City_
    ↳ Occurences")
```



## 2 Add Latitude and Longitude for Each City

We can now add the latitude and longitude for each city in the headlines. We will not add the country locations since there are a limited number of countries.

### 2.1 Accented Names

For finding the cities, we need to use accented city names. We'll create an accented name column using our mapping from the previous section.

```
[14]: import geonamescache
import unicode

import json

gc = geonamescache.GeonamesCache()

# Read in the saved unaccented:accented mapping
with open("../data/city_accent_mapping.json", "r") as fin:
    city_accented_mapping = json.loads(fin.read())

# Create a column for accented cities
data["accented_city"] = data["city"].map(city_accented_mapping)
data[data["city"] != data["accented_city"]].head()
```

```
[14]:
```

	headline	country	city	accented_city
7	Geneve Scientists Battle to Find Cure	NaN	Geneve	Genève
9	Zika Infested Monkeys in Sao Paulo	NaN	Sao Paulo	São Paulo
17	Louisiana Zika cases up to 26	NaN	NaN	NaN
19	Zika infects pregnant woman in Cebu	NaN	NaN	NaN
47	18 new Zika Cases in Bogota	NaN	Bogota	Bogotá

We can see there are several cases where the accented city does not match the original city.

```
[15]: print(gc.get_cities_by_name('São Paulo'))

[{'3448439': {'geonameid': 3448439, 'name': 'São Paulo', 'latitude': -23.5475,
'longitude': -46.63611, 'countrycode': 'BR', 'population': 10021295, 'timezone':
'America/Sao_Paulo', 'admin1code': '27'}}]
```

```
[16]: print(gc.get_cities_by_name('Sao Paulo'))
```

```
[]
```

We see the importance of using the accented names!

## 2.2 Handling Duplicate Cities

This is where we'll handle the duplicate cities. Our approach is relatively basic:

**For each city with multiple entries in geonames, we'll choose the city with the greatest population.**

This may occasionally be wrong, but a headline is more likely to mention a larger city (by population).

We can implement this by checking which is the largest entry for each city. Some cities have multiple locations as shown by Boston.

```
[17]: city = 'Boston'
gc.get_cities_by_name(city)
```

```
[17]: [{'2655138': {'geonameid': 2655138,
'name': 'Boston',
'latitude': 52.97633,
'longitude': -0.02664,
'countrycode': 'GB',
'population': 41340,
'timezone': 'Europe/London',
'admin1code': 'ENG'}},
{'4930956': {'geonameid': 4930956,
'name': 'Boston',
'latitude': 42.35843,
'longitude': -71.05977,
'countrycode': 'US',
```

```

'population': 667137,
'timezone': 'America/New_York',
'admin1code': 'MA'}}]

```

In this case we want Boston in the United States since it has the larger population. To get the largest city, we sort the matches by the `population` key.

```

[18]: matches = gc.get_cities_by_name(city)
matches = [{k: v for k, v in list(match.values())[0].items()} for match in matches]
matches = sorted(matches, key=lambda x: x["population"], reverse=True)
matches

```

```

[18]: [{'geonameid': 4930956,
      'name': 'Boston',
      'latitude': 42.35843,
      'longitude': -71.05977,
      'countrycode': 'US',
      'population': 667137,
      'timezone': 'America/New_York',
      'admin1code': 'MA'},
      {'geonameid': 2655138,
      'name': 'Boston',
      'latitude': 52.97633,
      'longitude': -0.02664,
      'countrycode': 'GB',
      'population': 41340,
      'timezone': 'Europe/London',
      'admin1code': 'ENG'}]

```

This sorts by the population of the cities descending (largest to smallest). If we take the first city, then we'll have the largest.

## 2.3 Finding Locations for Cities

Now let's find the locations of all the cities in the headlines. We'll want to be careful to go through the accented city names. If there are multiple matches for a city, we'll take the largest city.

```

[19]: city_locations = []

# Go through all the accented cities
for city in data["accented_city"]:
    # Find matches (if any)
    matches = gc.get_cities_by_name(city)
    if matches:
        # Sort from largest to smallest population
        matches = [

```



```

        {k: v for k, v in list(match.values())[0].items()} for match in
↪matches
    ]
    matches = sorted(matches, key=lambda x: x["population"], reverse=True)

    # Find the match with the largest population
    match = matches[0]

    # Record the information
    city_locations.append(
        {
            "name": match["name"],
            "latitude": match["latitude"],
            "longitude": match["longitude"],
            "countrycode": match["countrycode"],
            "pop": match["population"],
        }
    )

city_locations[-5:]

```

```

[19]: [{ 'name': 'Jerusalem',
        'latitude': 31.76904,
        'longitude': 35.21633,
        'countrycode': 'IL',
        'pop': 801000},
       { 'name': 'Indang',
        'latitude': 14.19528,
        'longitude': 120.87694,
        'countrycode': 'PH',
        'pop': 41159},
       { 'name': 'Suva',
        'latitude': -18.14161,
        'longitude': 178.44149,
        'countrycode': 'FJ',
        'pop': 77366},
       { 'name': 'Bella Vista',
        'latitude': 18.45539,
        'longitude': -69.9454,
        'countrycode': 'DO',
        'pop': 175683},
       { 'name': 'Wichita Falls',
        'latitude': 33.91371,
        'longitude': -98.49339,
        'countrycode': 'US',
        'pop': 104710}]

```

We can convert this list of dictionaries to a dataframe.

```
[20]: city_locations = pd.DataFrame(city_locations)
      city_locations.tail()
```

```
[20]:
```

	name	latitude	longitude	countrycode	pop
600	Jerusalem	31.76904	35.21633	IL	801000
601	Indang	14.19528	120.87694	PH	41159
602	Suva	-18.14161	178.44149	FJ	77366
603	Bella Vista	18.45539	-69.94540	DO	175683
604	Wichita Falls	33.91371	-98.49339	US	104710

```
[21]: city_locations = city_locations.drop_duplicates()
      print(f"We have the locations for {city_locations.shape[0]} unique cities.")
```

We have the locations for 573 unique cities.

Next let's merge with the headlines on the `accented_city` and `name`.

```
[22]: data = pd.merge(
      data, city_locations, left_on="accented_city", right_on="name", how="left"
    )
      data.head()
```

```
[22]:
```

	headline	country	city	\
0	Zika Outbreak Hits Miami	NaN	Miami	
1	Could Zika Reach New York City?	NaN	New York City	
2	First Case of Zika in Miami Beach	NaN	Miami Beach	
3	Mystery Virus Spreads in Recife, Brazil	Brazil	Recife	
4	Dallas man comes down with case of Zika	NaN	Dallas	

	accented_city	name	latitude	longitude	countrycode	pop
0	Miami	Miami	25.77427	-80.19366	US	441003.0
1	New York City	New York City	40.71427	-74.00597	US	8175133.0
2	Miami Beach	Miami Beach	25.79065	-80.13005	US	92312.0
3	Recife	Recife	-8.05389	-34.88111	BR	1478098.0
4	Dallas	Dallas	32.78306	-96.80667	US	1300092.0

Let's make sure keeping the largest city worked. We can try Boston as well as Rochester, both of which should be in the United States.

```
[23]: data[data['city'] == 'Boston']
```

```
[23]:
```

	headline	country	city	accented_city	name	latitude	\
27	Flu season hits Boston	NaN	Boston	Boston	Boston	42.35843	

	longitude	countrycode	pop
27	-71.05977	US	667137.0

```
[24]: data[data['city'] == 'Rochester']
```

```
[24]:
```

	headline	country	city	\
84	Rochester authorities confirmed the spread of ...	NaN	Rochester	
298	Herpes Keeps Spreading in Rochester	NaN	Rochester	


	accented_city	name	latitude	longitude	countrycode	pop
84	Rochester	Rochester	43.15478	-77.61556	US	209802.0
298	Rochester	Rochester	43.15478	-77.61556	US	209802.0

It looks like our method for finding the largest city worked as expected. For each headline with a city in `geonamescache`, we now have the latitude and longitude.

Due to the limited number of countries found in headlines, we'll stick to only the cities.

## 2.4 Data Cleaning

```
[25]: _ = data.isna().sum().plot.bar(title='Missing Values by Column')
```



We can see there are quite a few missing values in the `country` column. Let's just remove the country since it does not give us much information.

```
[26]: data = data.drop(columns=['country'])
```

Let's investigate the headlines where we don't have a `name`. We might be able to figure out more data cleaning steps to take.

```
[27]: pd.options.display.max_colwidth = 100

no_name = data[data["name"].isna()].copy()

print(f"There are {len(no_name)} headlines without a city.")

no_name.tail()
```

There are 42 headlines without a city.

```
[27]:
```

	headline	city	\
596	Zika arrives in Dangriga	NaN	
601	More Patients in Maynard are Getting Diagnosed with Syphilis	NaN	
625	Zika case reported in Antioquia	NaN	
627	Chikungunya has not Left Pismo Beach	NaN	
628	Zika spreads to La Joya	NaN	

	accented_city	name	latitude	longitude	countrycode	pop
596	NaN	NaN	NaN	NaN	NaN	NaN
601	NaN	NaN	NaN	NaN	NaN	NaN
625	NaN	NaN	NaN	NaN	NaN	NaN
627	NaN	NaN	NaN	NaN	NaN	NaN
628	NaN	NaN	NaN	NaN	NaN	NaN

We should manually check a few of these to make sure we can't find a city for the headline.

```
[28]: city_set = set(city_accented_mapping.keys())

for city in ["Dangriga", "Maynard", "Antioquia", "Pismo Beach", "La Joya"]:
    if city in city_set:
        print("Found ", city)
    else:
        print("Did Not Find City")
```

```
Did Not Find City
Did Not Find City
Did Not Find City
Did Not Find City
Did Not Find City
```

It appears that the 42 headlines without a city name may have a city, but it is not included in `geonamescache`. We'll have to go ahead and remove these cities since they cannot be used.

```
[29]: data = data.dropna(subset=['name'])
data.describe()
```


```
[29]:
```

	latitude	longitude	pop
count	605.000000	605.000000	6.050000e+02
mean	26.765746	-38.243197	8.904713e+05
std	20.619771	79.480854	1.974091e+06
min	-53.787690	-156.506040	1.338100e+04
25%	16.419040	-90.444300	5.878700e+04
50%	33.749000	-76.496610	1.712140e+05
75%	40.714270	7.095490	6.480340e+05
max	59.938630	179.364510	1.307630e+07

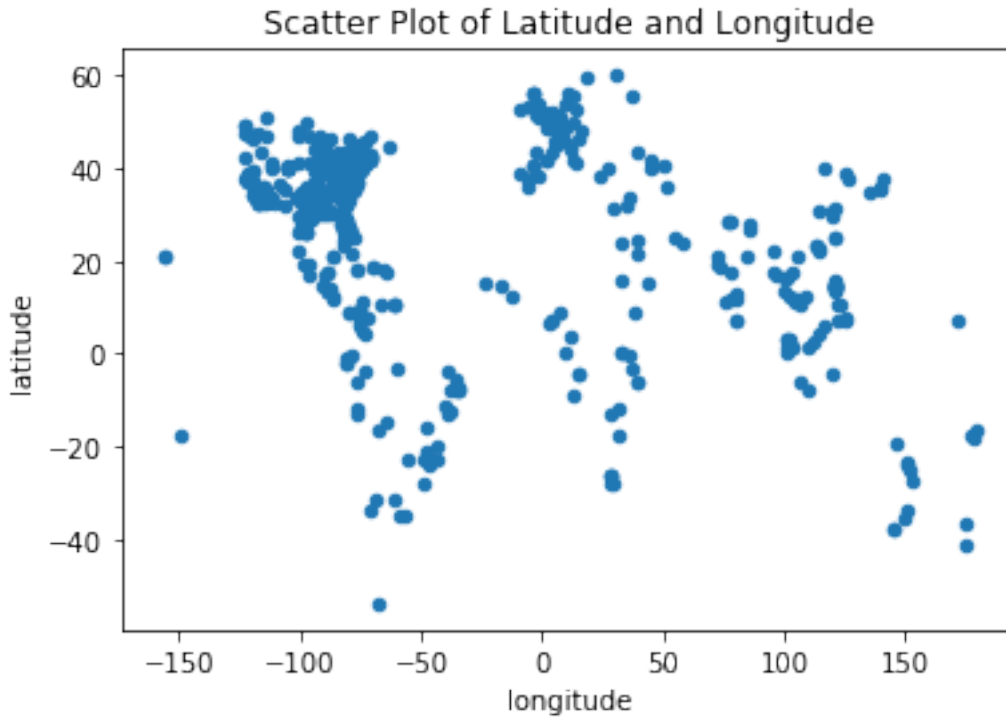
```
[30]: print(f'We have the city locations for {len(data)} cities.')
```

We have the city locations for 605 cities.

```
[31]: _ = data['pop'].plot.hist(title='Histogram of Population')
```



```
[32]: _ = data.plot.scatter(
    x="longitude", y="latitude", title="Scatter Plot of Latitude and Longitude"
)
```



This looks sort of like a map! We'll have to use a map in the next section.

As a final cleaning step, we can remove the `name` column since it is redundant with `city`. The final data frame is below.

```
[33]: data = data.drop(columns=['name'])
      data.tail()
```

```
[33]:
```

				headline \		
642	Rumors about Rabies spreading in Jerusalem have been refuted					
643	More Zika patients reported in Indang					
644	Suva authorities confirmed the spread of Rotavirus					
645	More Zika patients reported in Bella Vista					
646	Zika Outbreak in Wichita Falls					

	city	accented_city	latitude	longitude	countrycode	pop
642	Jerusalem	Jerusalem	31.76904	35.21633	IL	801000.0
643	Indang	Indang	14.19528	120.87694	PH	41159.0
644	Suva	Suva	-18.14161	178.44149	FJ	77366.0
645	Bella Vista	Bella Vista	18.45539	-69.94540	DO	175683.0
646	Wichita Falls	Wichita Falls	33.91371	-98.49339	US	104710.0

```
[34]: data[['headline', 'city', 'latitude', 'longitude', 'countrycode']].head(10)
```

```
[34]:
```

	headline	city	latitude	\
0	Zika Outbreak Hits Miami	Miami	25.77427	
1	Could Zika Reach New York City?	New York City	40.71427	
2	First Case of Zika in Miami Beach	Miami Beach	25.79065	
3	Mystery Virus Spreads in Recife, Brazil	Recife	-8.05389	
4	Dallas man comes down with case of Zika	Dallas	32.78306	
5	Trinidad confirms first Zika case	Trinidad	-14.83333	
6	Zika Concerns are Spreading in Houston	Houston	29.76328	
7	Geneve Scientists Battle to Find Cure	Geneve	46.20222	
8	The CDC in Atlanta is Growing Worried	Atlanta	33.74900	
9	Zika Infested Monkeys in Sao Paulo	Sao Paulo	-23.54750	

	longitude	countrycode
0	-80.19366	US
1	-74.00597	US
2	-80.13005	US
3	-34.88111	BR
4	-96.80667	US
5	-64.90000	BO
6	-95.36327	US
7	6.14569	CH
8	-84.38798	US
9	-46.63611	BR

This dataframe is the final outcome from this section. We will use it to cluster headlines based on the geographic location in the next section.

## 2.5 Saving Data

Let's save the final processed dataframe to a csv file for easy input and output with Pandas.

```
[35]: data.to_csv('../data/processed_headlines_locations.csv')
```

## 3 Summary

In this notebook we:

- Read the parsed headlines into a dataframe
- Found the location of the cities mentioned in the headlines
- Kept the largest city if a city was in geonames multiple times
- Joined the cities to the headlines
- Cleaned up the final dataframe to only headlines with a location

The end deliverable is a dataframe containing the headline, the city mentioned in the headline, the location of the city, and the population of the city. We can move on to clustering and visualizing the headline locations in the next section!

```
[ ]:
```